

# STAT\_MECH\_ASSIGNMENT\_02

September 10, 2025

## 1 Multivariate Gaussian Integral (n = 1..10) — Algorithm & Pseudocode with Explanations

### 1.1 Problem

We want to compute, for dimensions  $n = 1, 2, \dots, 10$ , the **multivariate Gaussian integral**:

$$I_n = \int_{\mathbb{R}^n} \exp\left(-\frac{1}{2}x^\top Ax + b^\top x\right) dx$$

Here:

- $x$  is a vector of length  $n$ ,
- $A$  is an  $n \times n$  **symmetric positive definite matrix**,
- $b$  is a vector of length  $n$ .

We will compare two ways of computing this integral:

1. The **analytic** value — using known formulas from statistics / linear algebra,
2. A **numerical** estimate — using **importance sampling**, a Monte Carlo method.

Finally, we will plot the **relative error** between the two approaches as a function of dimension  $n$ .

---

### 1.2 Inputs

To keep things simple and reproducible, we fix the following:

- **Dimension:**  $n \in \{1, 2, \dots, 10\}$  — the size of the vector and matrix
- **Matrix  $A$ :** A **tridiagonal matrix** with:
  - diagonal elements  $d_0 = 3.0$ ,
  - off-diagonal elements  $a = 0.4$ ,
  - which guarantees  $A$  is **symmetric positive definite** (i.e., integral converges).
- **Vector  $b$ :** A random vector drawn from a standard normal distribution  $\mathcal{N}(0, 1)^n$ . We use a fixed seed so the results are the same every time.
- **Monte Carlo samples  $N$ :** Number of random samples to estimate the integral numerically. We choose:

$$N = 60,000$$

- **Proposal distribution standard deviation  $s$ :** The importance sampler uses a Gaussian centered at the mode. We set:

$$s = \frac{1}{\sqrt{d_0}} \approx 0.577$$


---

### 1.3 Outputs

For each dimension  $n$ , we will record and report:

- The **analytic value** of the integral
- The **numeric estimate** (from sampling)
- The **standard error** (how noisy the numeric estimate is)
- The **relative error** (how close numeric is to analytic)

Then, we will **plot relative error vs. dimension  $n$** .

---

### 1.4 Algorithm Overview

#### 1.4.1 Step-by-step for each $n = 1, 2, \dots, 10$ :

---

##### 1.4.2 Step 1: Construct the matrix $A$

We create an  $n \times n$  matrix that is:

- **Symmetric:**  $A = A^\top$
- **Tridiagonal:** only diagonal and nearest neighbors are non-zero
- **Toeplitz:** same values along diagonals

It looks like this (for example, if  $n = 4$ ):

$$A = \begin{bmatrix} 3.0 & 0.4 & 0 & 0 \\ 0.4 & 3.0 & 0.4 & 0 \\ 0 & 0.4 & 3.0 & 0.4 \\ 0 & 0 & 0.4 & 3.0 \end{bmatrix}$$

This matrix has **fixed eigenvalues bounded away from zero**, so the integral is always well-defined.

---

### 1.4.3 Step 2: Draw a random vector $b$

We generate a random vector  $b \in \mathbb{R}^n$ , where each element is drawn independently from a standard normal distribution  $\mathcal{N}(0, 1)$ .

This vector gives us a linear term in the exponent:  $b^\top x$ .

---

### 1.4.4 Step 3: Compute the analytic value

We use the closed-form formula for the integral of a multivariate Gaussian of the form:

$$\int \exp\left(-\frac{1}{2}x^\top Ax + b^\top x\right) dx = \exp\left(\frac{n}{2}\log(2\pi) - \frac{1}{2}\log\det A + \frac{1}{2}b^\top A^{-1}b\right)$$

Here's how we compute it:

- Solve  $A\mu = b$  to get  $\mu = A^{-1}b$  (we never compute  $A^{-1}$  directly — just solve the equation)
- Compute the **quadratic term**:  $b^\top \mu$
- Compute  $\log\det A$  using a stable method (`np.linalg.slogdet`)
- Plug into the formula:

$$\log I = \frac{n}{2}\log(2\pi) - \frac{1}{2}\log\det A + \frac{1}{2}b^\top \mu$$

- Finally, exponentiate to get:

$$I_{\text{analytic}} = \exp(\log I)$$

---

### 1.4.5 Step 4: Estimate numerically using importance sampling

We now approximate the integral using **importance sampling** which is Monte Carlo technique.

We sample from a Gaussian proposal distribution:

$$q(x) = \mathcal{N}(\mu, s^2 I)$$

This distribution is centered at the mode of the integrand (i.e., at  $\mu = A^{-1}b$ ), and has variance  $s^2$  in every direction.

---

**For each sample  $i = 1$  to  $N$ :**

1. **Draw sample:**

$$x_i = \mu + sz_i \quad \text{where } z_i \sim \mathcal{N}(0, I)$$

2. **Evaluate the integrand** at  $x_i$ :

$$\log f_i = -\frac{1}{2}x_i^\top A x_i + b^\top x_i$$

3. **Evaluate the proposal density**  $q$  at  $x_i$ :

$$\log q_i = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(s^2) - \frac{1}{2s^2} \|x_i - \mu\|^2$$

4. **Compute importance weight**:

$$w_i = \exp(\log f_i - \log q_i)$$

*(In practice, we use the log-sum-exp trick to avoid underflow.)*

---

**After all  $N$  samples:**

- Compute the estimate of the integral:

$$I_{\text{numeric}} = \frac{1}{N} \sum_{i=1}^N w_i$$

- Compute the standard error:

$$\text{SE} = \sqrt{\left( \frac{1}{N} \sum_{i=1}^N w_i^2 - I_{\text{numeric}}^2 \right) / N}$$

This gives us a numerical estimate **with error bars**.

---

#### 1.4.6 Step 5: Compute relative error

Compare the numerical estimate with the true value:

$$\text{RelErr} = \frac{|I_{\text{numeric}} - I_{\text{analytic}}|}{I_{\text{analytic}}}$$

A small relative error means the Monte Carlo method is accurate.

---

## 1.5 Final Step: Plot

We plot the **dimension**  $n$  on the x-axis and **relative error** on the y-axis.

This shows how the quality of the numerical estimate behaves as we increase the dimension.

## 1.6 Summary

Step	Description
Build $A$	Fixed-diagonal, tridiagonal Toeplitz matrix
Draw $b$	From standard normal
Compute analytic $I$	Use known formula from Gaussian integrals
Estimate numeric $I$	Using importance sampling with Gaussian proposal
Compute error	Relative error between numeric and analytic values
Plot	Error vs dimension

```
[3]: import numpy as np
import matplotlib.pyplot as plt

def make_tridiagonal_toeplitz(n, diag_value=3.0, offdiag_value=0.4):
    A = np.zeros((n, n), dtype=float)
    np.fill_diagonal(A, diag_value)
    if n > 1:
        A[np.arange(n-1), np.arange(1, n)] = offdiag_value
        A[np.arange(1, n), np.arange(n-1)] = offdiag_value
    return A

def gaussian_integral_analytic(A, b):
    n = A.shape[0]
    mu = np.linalg.solve(A, b)
    quad = float(b @ mu)
    sign, logdetA = np.linalg.slogdet(A)
    if sign <= 0:
        raise ValueError("Matrix A must have positive determinant.")
    logI = 0.5 * n * np.log(2 * np.pi) - 0.5 * logdetA + 0.5 * quad
    I_value = float(np.exp(logI))
    return I_value, logI, mu

def importance_sampling_isotropic(A, b, mu, scale, N_samples, rng):
    n = A.shape[0]
    s = float(scale)
    s2 = s * s
    batch = min(20000, N_samples)
    total = 0
    sum_w = 0.0
    sum_w2 = 0.0
```

```

while total < N_samples:
    k = min(batch, N_samples - total)
    Z = rng.standard_normal((k, n))
    X = mu + s * Z
    AX = X @ A.T
    logf = -0.5 * np.sum(X * AX, axis=1) + (X @ b)
    diff = X - mu
    quad_q = 0.5 * (1.0 / s2) * np.sum(diff * diff, axis=1)
    logq = -0.5 * n * np.log(2 * np.pi) - 0.5 * n * np.log(s2) - quad_q
    logw = logf - logq
    m = np.max(logw)
    w = np.exp(logw - m)
    sum_w += w.sum() * np.exp(m)
    sum_w2 += np.sum(w * w) * np.exp(2 * m)
    total += k
I_hat = sum_w / N_samples
var_hat = max(sum_w2 / N_samples - I_hat * I_hat, 0.0)
se = np.sqrt(var_hat / N_samples)
return I_hat, se

def main():
    diag_value = 3.0
    offdiag_value = 0.4
    base_seed = 2025
    rng_global = np.random.default_rng(base_seed)
    proposal_scale = 1.0 / np.sqrt(diag_value)
    N_samples = 60000
    ns = list(range(1, 101))
    I_true_list = []
    I_num_list = []
    SE_list = []
    rel_err_list = []
    for n in ns:
        A = make_tridiagonal_toeplitz(n, diag_value=diag_value,
        ↪offdiag_value=offdiag_value)
        b = rng_global.normal(size=n)
        I_true, logI, mu = gaussian_integral_analytic(A, b)
        rng_is = np.random.default_rng(base_seed + n)
        I_hat, se = importance_sampling_isotropic(
            A=A,
            b=b,
            mu=mu,
            scale=proposal_scale,
            N_samples=N_samples,
            rng=rng_is
        )
        rel_err = abs(I_hat - I_true) / I_true

```

```

I_true_list.append(I_true)
I_num_list.append(I_hat)
SE_list.append(se)
rel_err_list.append(rel_err)

print(" n |      I_analytic      I_numeric(IS)      SE(IS)      \
↪RelErr")
for n, Ia, In, se, re in zip(ns, I_true_list, I_num_list, SE_list, \
↪rel_err_list):
    print(f"{n:2d} | {Ia:14.6e}  {In:14.6e}  {se:9.2e}  {re:8.3e}")

plt.figure()
plt.plot(ns, rel_err_list, marker='o')
plt.xlabel("Dimension n")
plt.ylabel("Relative error (numeric vs analytic)")
plt.title("Multivariate Gaussian integral - relative error vs dimension (n \
↪= 1..100)")
plt.grid(True)
plt.show()

main()


```

n	I_analytic	I_numeric(IS)	SE(IS)	RelErr
1	3.293532e+00	3.293532e+00	0.00e+00	2.697e-16
2	2.221701e+00	2.221857e+00	1.24e-03	7.020e-05
3	1.132531e+01	1.132175e+01	9.33e-03	3.142e-04
4	5.496558e+00	5.496008e+00	5.72e-03	9.994e-05
5	1.289030e+01	1.289922e+01	1.58e-02	6.921e-04
6	1.528187e+02	1.533375e+02	2.15e-01	3.395e-03
7	5.362022e+01	5.356167e+01	7.98e-02	1.092e-03
8	2.522477e+01	2.529694e+01	4.22e-02	2.861e-03
9	1.759672e+02	1.758965e+02	3.08e-01	4.016e-04
10	3.328045e+03	3.326363e+03	6.33e+00	5.052e-04
11	3.227120e+02	3.224488e+02	6.45e-01	8.156e-04
12	5.082845e+02	5.096214e+02	1.07e+00	2.630e-03
13	2.227162e+03	2.233761e+03	4.96e+00	2.963e-03
14	2.403444e+04	2.413115e+04	5.79e+01	4.024e-03
15	4.934698e+03	4.929799e+03	1.17e+01	9.929e-04
16	9.512448e+03	9.532346e+03	2.49e+01	2.092e-03
17	4.340738e+04	4.351096e+04	1.13e+02	2.386e-03
18	1.593466e+04	1.595442e+04	4.30e+01	1.240e-03
19	4.177025e+04	4.167958e+04	1.17e+02	2.171e-03
20	3.438880e+05	3.442915e+05	1.01e+03	1.173e-03
21	1.504419e+05	1.506268e+05	4.42e+02	1.229e-03
22	1.866835e+05	1.865809e+05	5.75e+02	5.495e-04
23	5.917952e+05	5.902660e+05	1.84e+03	2.584e-03
24	2.218064e+05	2.218009e+05	7.08e+02	2.476e-05

25	8.585698e+04	8.570794e+04	2.97e+02	1.736e-03
26	1.162908e+06	1.166407e+06	4.57e+03	3.008e-03
27	2.660043e+06	2.669138e+06	9.70e+03	3.419e-03
28	2.409291e+07	2.405072e+07	8.66e+04	1.751e-03
29	3.283429e+06	3.279472e+06	1.20e+04	1.205e-03
30	2.812848e+06	2.810516e+06	1.08e+04	8.291e-04
31	1.121551e+07	1.125404e+07	4.31e+04	3.435e-03
32	2.952011e+08	2.971417e+08	1.18e+06	6.574e-03
33	1.251703e+08	1.253791e+08	5.50e+05	1.668e-03
34	7.659040e+07	7.622551e+07	3.32e+05	4.764e-03
35	1.534217e+08	1.531034e+08	6.34e+05	2.074e-03
36	3.181361e+08	3.157562e+08	1.34e+06	7.481e-03
37	2.215166e+08	2.199724e+08	9.31e+05	6.971e-03
38	4.045977e+08	4.046273e+08	1.98e+06	7.310e-05
39	3.287243e+08	3.280319e+08	1.49e+06	2.107e-03
40	1.227328e+09	1.237686e+09	6.08e+06	8.440e-03
41	4.546773e+08	4.558403e+08	2.21e+06	2.558e-03
42	9.552569e+10	9.558925e+10	4.64e+08	6.654e-04
43	6.876634e+10	6.824370e+10	3.15e+08	7.600e-03
44	2.537347e+10	2.529055e+10	1.24e+08	3.268e-03
45	7.648426e+09	7.625288e+09	3.91e+07	3.025e-03
46	1.618550e+12	1.612512e+12	8.70e+09	3.731e-03
47	1.092450e+11	1.091056e+11	5.87e+08	1.277e-03
48	1.406911e+11	1.417207e+11	7.74e+08	7.318e-03
49	7.995554e+10	7.990225e+10	4.33e+08	6.665e-04
50	2.893612e+11	2.903569e+11	1.59e+09	3.441e-03
51	1.413259e+12	1.414023e+12	8.20e+09	5.411e-04
52	1.916243e+12	1.912535e+12	1.19e+10	1.935e-03
53	6.694349e+11	6.693081e+11	3.87e+09	1.893e-04
54	2.888176e+12	2.870242e+12	1.67e+10	6.210e-03
55	6.467820e+12	6.540971e+12	4.35e+10	1.131e-02
56	2.462296e+14	2.457885e+14	1.55e+12	1.791e-03
57	3.320469e+14	3.317829e+14	2.19e+12	7.951e-04
58	8.644614e+14	8.590514e+14	5.21e+12	6.258e-03
59	3.139316e+13	3.139180e+13	1.98e+11	4.328e-05
60	3.704676e+13	3.716419e+13	2.38e+11	3.170e-03
61	7.980005e+15	7.934670e+15	5.05e+13	5.681e-03
62	1.927169e+15	1.894894e+15	1.19e+13	1.675e-02
63	3.818057e+15	3.813720e+15	2.72e+13	1.136e-03
64	8.562324e+15	8.529402e+15	5.64e+13	3.845e-03
65	5.011401e+14	5.040481e+14	3.51e+12	5.803e-03
66	4.668067e+15	4.679648e+15	3.54e+13	2.481e-03
67	4.721915e+15	4.698978e+15	3.26e+13	4.858e-03
68	1.299731e+16	1.291734e+16	8.92e+13	6.152e-03
69	1.271445e+18	1.255676e+18	8.78e+15	1.240e-02
70	4.529980e+17	4.538845e+17	3.17e+15	1.957e-03
71	9.224145e+17	9.199861e+17	7.02e+15	2.633e-03
72	5.722506e+17	5.757161e+17	4.47e+15	6.056e-03



73	5.067122e+17	5.036112e+17	3.68e+15	6.120e-03
74	2.547230e+18	2.578668e+18	2.05e+16	1.234e-02
75	4.452358e+17	4.436285e+17	3.38e+15	3.610e-03
76	1.830800e+18	1.827620e+18	1.43e+16	1.737e-03
77	2.528663e+18	2.542209e+18	1.96e+16	5.357e-03
78	4.916899e+17	4.928271e+17	4.13e+15	2.313e-03
79	4.598401e+18	4.672194e+18	4.10e+16	1.605e-02
80	1.149307e+18	1.129323e+18	8.54e+15	1.739e-02
81	1.643041e+20	1.648743e+20	1.43e+18	3.471e-03
82	1.500203e+19	1.474679e+19	1.17e+17	1.701e-02
83	3.605695e+17	3.608619e+17	3.79e+15	8.110e-04
84	1.234446e+19	1.233520e+19	1.24e+17	7.503e-04
85	3.265917e+20	3.274765e+20	2.87e+18	2.709e-03
86	1.257650e+20	1.253338e+20	1.19e+18	3.429e-03
87	9.734068e+19	9.820293e+19	9.11e+17	8.858e-03
88	4.986568e+20	4.989637e+20	4.70e+18	6.155e-04
89	5.604125e+20	5.656556e+20	5.07e+18	9.356e-03
90	1.376800e+22	1.397188e+22	1.51e+20	1.481e-02
91	5.755367e+21	5.849657e+21	5.74e+19	1.638e-02
92	1.501075e+22	1.484111e+22	1.37e+20	1.130e-02
93	8.523204e+22	8.608627e+22	8.44e+20	1.002e-02
94	1.051436e+22	1.037280e+22	9.52e+19	1.346e-02
95	3.761416e+22	3.754465e+22	3.64e+20	1.848e-03
96	9.439982e+22	9.367638e+22	9.03e+20	7.664e-03
97	1.562788e+21	1.537862e+21	1.41e+19	1.595e-02
98	4.940941e+25	4.884686e+25	5.22e+23	1.139e-02
99	1.061939e+25	1.074427e+25	1.21e+23	1.176e-02
100	8.262764e+22	8.301455e+22	8.35e+20	4.683e-03



```
[8]: import numpy as np
from math import sqrt

def pi(N, seed=12345):
    rng = np.random.default_rng(seed)

    # Generate N pairs (x, y) uniformly in [-1, 1] x [-1, 1]
    xy = rng.uniform(-1.0, 1.0, size=(N, 2))

    # Count how many fall inside the unit circle
    in_circle = np.count_nonzero(np.sum(xy**2, axis=1) <= 1.0)

    # Estimate using 4 * (points in circle / total points)
    p_hat = in_circle / N
    pi_hat = 4.0 * p_hat

    # Standard error: SE = 4 * sqrt(p(1-p)/N)
    se = 4.0 * sqrt(p_hat * (1.0 - p_hat) / N)

    # 95% confidence interval
    ci_lower = pi_hat - 1.96 * se
    ci_upper = pi_hat + 1.96 * se
```

```

    return {
        "N": N,
        "In circle": in_circle,
        " estimate": pi_hat,
        "Standard error ( $\pm 1$ )": se,
        "95% CI": (ci_lower, ci_upper)
    }

if __name__ == "__main__":
    trials = [10_000, 100_000, 1_000_000, 2_000_000]
    for N in trials:
        result = pi(N)
        print(f"N = {result['N']:,}")
        print(f"   In circle: {result['In circle']}")
        print(f"       {result[' estimate']:.6f}  $\pm$  {result['Standard error_
 $\pm 1$ ']:.6f}")
        print(f"   95% CI: ({result['95% CI'][0]:.6f}, {result['95% CI'][1]:.
 $\pm 6f$ })\n")

```

```

N = 10,000
  In circle: 7883
    3.153200  $\pm$  0.016341
  95% CI: (3.121173, 3.185227)

```

```

N = 100,000
  In circle: 78435
    3.137400  $\pm$  0.005202
  95% CI: (3.127204, 3.147596)

```

```

N = 1,000,000
  In circle: 784595
    3.138380  $\pm$  0.001644
  95% CI: (3.135157, 3.141603)

```

```

N = 2,000,000
  In circle: 1569055
    3.138110  $\pm$  0.001163
  95% CI: (3.135831, 3.140389)

```

```

[11]: import numpy as np
import matplotlib.pyplot as plt

def pi_convergence(Nmax=5000, seed=123):
    rng = np.random.default_rng(seed)
    x = rng.uniform(-1.0, 1.0, size=Nmax)

```

```

y = rng.uniform(-1.0, 1.0, size=Nmax)
r2 = x**2 + y**2
inside = r2 <= 1.0

pi_estimates = np.cumsum(inside) / np.arange(1, Nmax + 1) * 4
true_pi = np.pi
error = np.abs(pi_estimates - true_pi)

fig, ax = plt.subplots(2, 1, figsize=(8, 6), sharex=True)

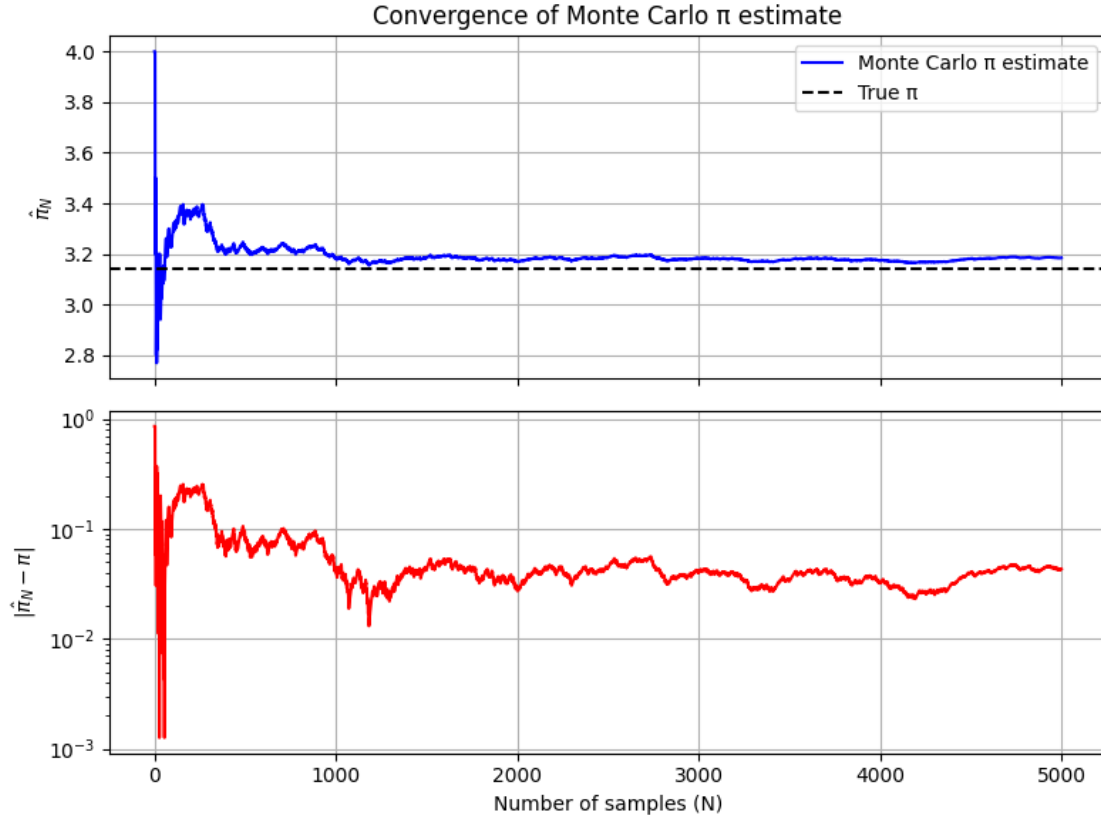
ax[0].plot(pi_estimates, label=' estimate', color='blue')
ax[0].axhline(true_pi, color='black', linestyle='--', label='True ')
ax[0].set_ylabel(r'$\hat{\pi}_N$')
ax[0].set_title("Convergence of estimate")
ax[0].legend()
ax[0].grid(True)

ax[1].plot(error, label='Absolute error', color='red')
ax[1].set_ylabel(r'$|\hat{\pi}_N - \pi|$')
ax[1].set_xlabel("Number of samples (N)")
ax[1].set_yscale("log")
ax[1].grid(True)

plt.tight_layout()
plt.show()

plot_mc_pi_convergence()

```



### 1.6.1 Estimation of $\pi$ Using Truncated Normal Sampling

In this method, we estimate the value of  $\pi$  using **Monte Carlo sampling**, but instead of drawing uniformly from the square  $[-1, 1]^2$ , we sample points from a **normal distribution**:

$$x, y \sim \mathcal{N}(0, \sigma^2), \quad \text{with } \sigma = 0.5$$

We then **accept only those points** that fall within the square  $[-1, 1]^2$ , i.e., we apply **rejection sampling**.

Unlike uniform sampling, where all regions of the square are equally likely, the normal distribution clusters points **closer to the origin**  $(0, 0)$ , which lies **inside the unit circle**. This creates a non-uniform sampling density that affects:

- **More points fall inside the circle early on**, improving initial convergence.
- However, **fewer points sample the edges** of the square, making the estimator **more biased** if not enough points are used.
- Overall convergence to  $\pi$  is still guaranteed by the **law of large numbers**, but it may **converge more slowly** than uniform sampling due to this skewed distribution.
- This method still estimates  $\pi$  by:

$$\hat{\pi}_N = 4 \cdot \frac{\text{\#points inside circle}}{\text{\#total accepted points}}$$

- But the accuracy and convergence depend on how the sampling distribution populates the square.
- Normal distributions **require more samples** to compensate for non-uniformity at the boundaries of the square.

[ ]: